

1 Supplementary materials for: Modelling
2 palaeoecological time series using
3 generalized additive models

4 *Gavin L. Simpson*

5 *August 15, 2018*

6 **1 Introduction**

- 7 This document is an annotated version of the R code used to fit the GAMs and related analyses
8 to the Small Water and Braya-Sø example data sets.
9 The following packages are required: *mgcv*, *scam*, *ggplot2*, *cowplot*, and *tidyverse*.

```
library("mgcv")
```

```
10 #> Loading required package: nlme
```

```
11 #> This is mgcv 1.8-24. For overview type 'help("mgcv-package")'.
```

```
library("scam")
```

```
12 #> This is scam 1.2-2.
```

```
library("ggplot2")
```

```
library("cowplot")
```

```
13 #>
```

```
14 #> Attaching package: 'cowplot'
```

```
15 #> The following object is masked from 'package:ggplot2':
```

```
16 #>
```

```
17 #>     ggsave
```

```
library("tidyverse")
```

- 18 In addition, the *schoenberg* package is required; this in-development package is not on CRAN
19 but can be installed directly from GitHub using functions from the *devtools* package. To install
20 the package, install *devtools* and then use *devtools::install_github()* to install *schoenberg*,
21 as shown below:

```
## schoenberg is not on CRAN, install from github:  
install.packages("devtools")  
devtools::install_github("gavinsimpson/schoenberg")
```

22 Assuming that the installation of *schoenberg* completes without error, the package can be loaded
23 as usual

```
library("schoenberg") # need to change the name of the package
```

24 The final environment-preparation step is to set the default *ggplot* theme, which the loading of
25 *cowplot* has over-ridden. Here I use the more-minimal black-and-white theme (*theme_bw()*)

```
## Default ggplot theme
theme_set(theme_bw())
```

2 Load the data sets

27 The example data sets are also stored on GitHub; <https://github.com/gavinsimpson/>
28 *frontiers-palaeo-additive-modelling*. Once downloaded the data are read in and processed a
29 little

```
## source Small Water data
small <- readRDS("./data/small-water/small-water-isotope-data.rds")
head(small)

#>   Depth   d13C TotalC d15N TotalN DryWeight      Year
#> 1  0.2 -27.57  806.49 3.05  64.21      8.2 2007.254
#> 2  0.4 -27.67  949.33 3.01  73.26      7.6 2006.510
#> 3  0.8 -27.63 1305.52 2.93  93.25     11.6 2004.941
#> 4  1.2 -27.62 1136.04 2.33  86.09      9.6 2003.269
#> 5  1.6 -27.48 1028.27 2.09  93.80     10.9 2001.496
#> 6  2.0 -27.39  809.91 2.66  79.98      9.9 1999.626

## load braya so data set
braya <- read.table("./data/braya-so/DAndrea.2011.Lake Braya So.txt",
                     skip = 84)
## clean up variable names
names(braya) <- c("Depth", "DepthUpper", "DepthLower", "Year", "YearYoung",
                  "YearOld", "UK37")
## add a variable for the amount of time per sediment sample
braya <- transform(braya, sampleInterval = YearYoung - YearOld)
head(braya)
```

```
#>   Depth DepthUpper DepthLower      Year YearYoung YearOld    UK37
#> 1  0.25          0.0        0.5 1999.125  2006.00 1992.25 -0.640
#> 2  0.75          0.5        1.0 1985.375  1992.25 1978.50 -0.637
#> 3  1.25          1.0        1.5 1971.525  1978.50 1964.55 -0.614
#> 4  1.75          1.5        2.0 1957.575  1964.55 1950.60 -0.627
#> 5  2.25          2.0        2.5 1943.150  1950.60 1935.70 -0.633
#> 6  2.75          2.5        3.0 1928.250  1935.70 1920.80 -0.616
```

```

44 #> sampleInterval
45 #> 1           13.75
46 #> 2           13.75
47 #> 3           13.95
48 #> 4           13.95
49 #> 5           14.90
50 #> 6           14.90

## plot labels
d15n_label <- expression(delta^{15}N)
braya_ylabel <- expression(italic(U)[37]^{\it k})

```

51 Plots of the two data sets are prepared using `ggplot2`

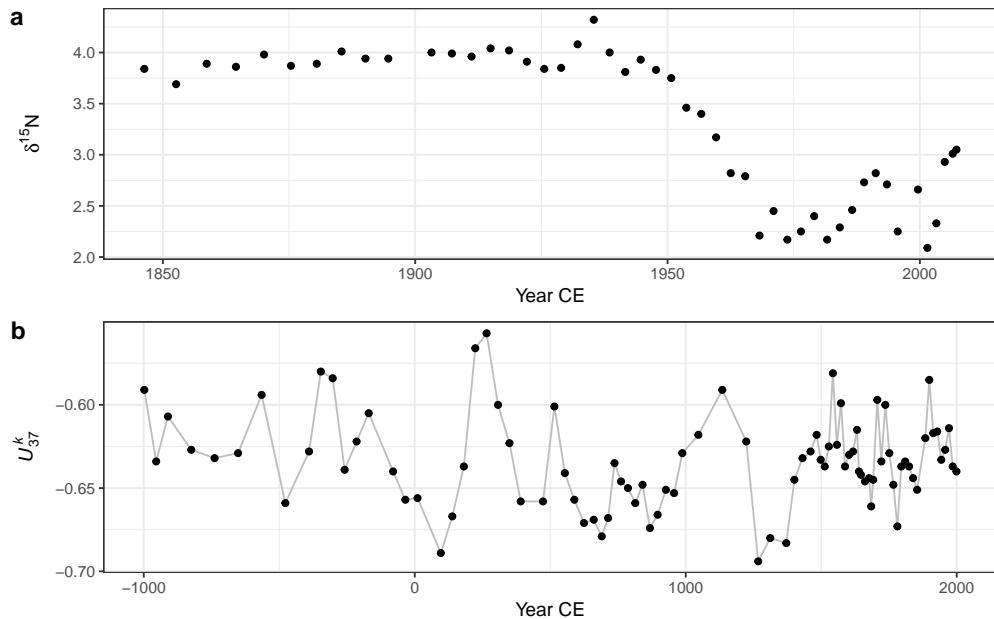
```

## plot Small Water data
small_plt <- ggplot(small, aes(x = Year, y = d15N)) +
  geom_point() +
  labs(y = d15n_label, x = "Year CE")

## Generate a plot of the data
braya_plt <- ggplot(braya, aes(x = Year, y = UK37)) +
  geom_line(colour = "grey") +
  geom_point() +
  labs(y = braya_ylabel, x = "Year CE")

## Recreate plot from manuscript
plot_grid(small_plt, braya_plt, ncol = 1, labels = "auto", align = "hv",
          axis = "lr")

```



52

53 3 Fitting GAMs

54 Instead of following the structure of the paper exactly, the subsequent sections focus on the
55 two example data sets in turn, beginning with Small Water.

56 3.1 Small Water

57 A GAM is typically fitted using the `gam()` function from the *mgcv* package. The typical call
58 includes a formula describing the response variable and the linear predictor separated by the
59 `~` (tilde) symbol. The linear predictor part of the model contains the smooth function of the
60 time variable in the data set, and is indicated using the `s()` function. Unless other options are
61 needed, we only need specify where the variables in the formula can be found via the `data`
62 argument, and that we wish to use REML smoothness selection.

63 Putting this together we have the following function call to fit a simple GAM to the Small Water
64 isotope data

```
m <- gam(d15N ~ s(Year, k = 15), data = small, method = "REML")
```

65 As discussed in the paper, this model assumes that residuals¹ are independent. To account
66 for residual temporal autocorrelation we can include a continuous-time first-order autoregres-
67 sive (CAR(1)) process in the model residuals for GAMs fitted to data that are conditionally
68 distributed Gaussian.

69 The GAM plus CAR(1) process is fitted to the Small Water data set using the `gamm()` function.
70 This fits GAMs as mixed effects models via the *nlme* package, which allows the use of corre-
71 lation structures in the model residuals via the `correlation` argument. Here, the `corCAR1()`
72 function is used to select the CAR(1) process and we specify the ordering of samples via the
73 `Year` variable in `small`.

```
## fit small water GAM using gamm() with a CAR(1)
mod <- gamm(d15N ~ s(Year, k = 15), data = small,
             correlation = corCAR1(form = ~ Year), method = "REML")
```

74 The estimated value of ϕ for the CAR(1) can be extracted from the fitted model via the `$lme`
75 component. Here we just extract the correlation structure component.

```
## estimate of phi and confidence interval
smallPhi <- intervals(mod$lme, which = "var-cov")$corStruct
smallPhi
```

```
76 #>           lower      est.      upper
77 #> Phi 0.2811107 0.6026967 0.8547543
78 #> attr(,"label")
79 #> [1] "Correlation structure:"
```

¹or equivalently that the observations, conditional upon the model, are independent.

80 The object returned by `gamm()` includes both the linear mixed model and GAM faces of the
81 model, and as a result we need to access the separate elements (`lme` and `gam` respectively)
82 when proceeding to explore the model fit. The model summary is prepared from the `$gam`
83 component of the fitted model

```
## summary object
summary(mod$gam)

#>
#> Family: gaussian
#> Link function: identity
#>
#> Formula:
#> d15N ~ s(Year, k = 15)
#>
#> Parametric coefficients:
#>             Estimate Std. Error t value Pr(>|t|)
#> (Intercept) 3.30909   0.03489   94.84 <2e-16 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Approximate significance of smooth terms:
#>             edf Ref.df      F p-value
#> s(Year) 7.954 7.954 47.44 <2e-16 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> R-sq.(adj) =  0.929
#> Scale est. = 0.037268 n = 48
```

105 The output shows the estimated complexity of the fitted smooth, expressed in terms of the ef-
106 fective degrees of freedom of the spline. An associated F statistic and test of the null hypoth-
107 esis of no trend (effect) are also shown. Here the estimated trend provides strong evidence
108 against this null.

109 The CAR(1) process plotted in Figure 10 of the manuscript was prepared using


```
## plot CAR(1) process
S <- seq(0, 50, length = 100)
car1 <- setNames(as.data.frame(t(outer(smallPhi, S, FUN = `^`)[1, , ])),
                 c("Lower", "Correlation", "Upper"))
car1 <- transform(car1, S = S)

car1Plt <- ggplot(car1, aes(x = S, y = Correlation)) +
  geom_ribbon(aes(ymax = Upper, ymin = Lower),
              fill = "black", alpha = 0.2) +
  geom_line() +
```

```

    ylab(expression(italic(h) * (list(Delta[t], varphi)))) +
    xlab(expression(Delta[t] ~ (years)))
car1Plt

```



110

111 The exponential decline in correlation with increasing separation is evident here; once samples
112 are ~ 10 years apart, there is little estimated dependence between them.

113 The next code chunk prepares a plot of the fitted GAMS. The general idea is to predict from
114 the fitted model for a fine grid of points over the range of the time variable. Below I plot the
115 trend for Small Water with an approximate 95% point-wise confidence interval that assumes
116 asymptotic normality

```

N <- 300      # number of points at which to evaluate the smooth

## create new data to predict at; 200 evenly-spaced values over `Year`
newYear <- with(small, data.frame(Year = seq(min(Year), max(Year),
                                             length.out = 200)))

## Predict from the fitted model; note we predict from the $gam part
newYear <- cbind(newYear,
                  data.frame(predict(mod$gam, newYear, se.fit = TRUE)))

## Create the confidence interval
crit.t <- qt(0.975, df = df.residual(mod$gam))
newYear <- transform(newYear,
                      upper = fit + (crit.t * se.fit),
                      lower = fit - (crit.t * se.fit))

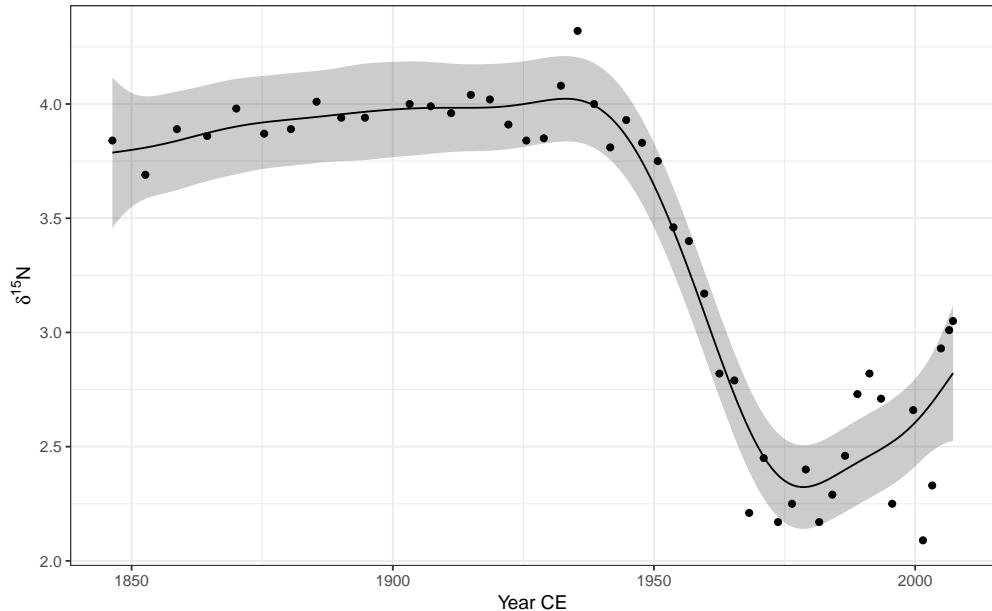
## Plot estimated trend
small_fitted <- ggplot(newYear, aes(x = Year, y = fit)) +
  geom_ribbon(aes(ymin = lower, ymax = upper, x = Year), alpha = 0.2,

```

```

            inherit.aes = FALSE, fill = "black") +
geom_point(data = small, mapping = aes(x = Year, y = d15N),
            inherit.aes = FALSE) +
geom_line() +
labs(y = d15n_label, x = "Year CE")
small_fitted

```



117

118 3.2 Braya-Sø

119 The same model is fitted to the Braya-Sø data set. Note however that in order to even fit the
 120 model with both a smooth and the CAR(1) process, I have had to change the default optimiser
 121 used to estimate the model parameters, and reduce the basis dimension to a small number.
 122 We also fit the model using GCV, which is the default, hence no method argument

```

## fit the car(1) model --- needs optim as this is not a stable fit!
## also needs k setting lower than default
braya.car1 <- gamm(UK37 ~ s(Year, k = 5), data = braya,
                     correlation = corCAR1(form = ~ Year),
                     method = "REML",
                     control = list(niterEM = 0, optimMethod = "BFGS",
                                   opt = "optim"))

## fit model using GCV
braya.gcv <- gam(UK37 ~ s(Year, k = 30), data = braya)

## estimate of phi and confidence interval
brayaPhi <- intervals(braya.car1$lme)$corStruct

```

```
brayaPhi
```

```
123 #>           lower      est. upper  
124 #> Phi 1.611049e-22 0.2000156      1  
125 #> attr(,"label")  
126 #> [1] "Correlation structure:"
```

127 Note the wide confidence interval — effectively 0–1 — on ϕ . If you were to increase the value of
128 k to be k = 10 in the s(Year) above, the model will fit but a warning message will be emitted
129 when trying to extract ϕ due to a non-positive definite model covariance matrix, indicating
130 problems with the model.

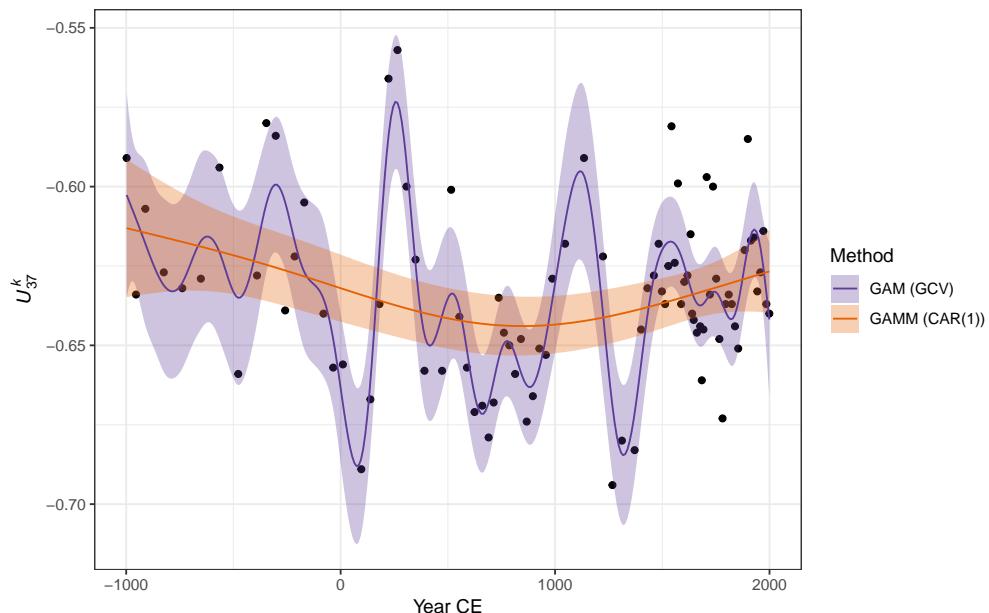
131 For Braya-Sø, we repeat the process, but we do so for both models (GAMM + CAR(1) and
132 GCV), and use a critical value from the t distribution to form the confidence interval

```
N <- 300    # number of points at which to evaluate the smooth  
  
## data to predict at  
newBraya <- with(braya, data.frame(Year = seq(min(Year), max(Year),  
                                         length.out = N)))  
## add predictions from GAMM + CAR(1) model  
newBraya <- cbind(newBraya,  
                    data.frame(predict(braya.car1$gam, newBraya,  
                                      se.fit = TRUE)))  
crit.t <- qt(0.975, df = df.residual(braya.car1$gam))  
newBraya <- transform(newBraya,  
                        upper = fit + (crit.t * se.fit),  
                        lower = fit - (crit.t * se.fit))  
  
## add GAM GCV results  
fit_gcv <- predict(braya.gcv, newdata = newBraya, se.fit = TRUE)  
crit.t <- qt(0.975, df.residual(braya.gcv))  
newGCV <- data.frame(Year = newBraya[["Year"]],  
                      fit = fit_gcv$fit,  
                      se.fit = fit_gcv$se.fit)  
newGCV <- transform(newGCV,  
                        upper = fit + (crit.t * se.fit),  
                        lower = fit - (crit.t * se.fit))  
newBraya <- rbind(newBraya, newGCV)          # bind on GCV results  
## Add indicator variable for model  
newBraya <- transform(newBraya,  
                        Method = rep(c("GAMM (CAR(1))", "GAM (GCV)"),  
                                      each = N))  
  
## plot CAR(1) and GCV fits  
braya_fitted <- ggplot(braya, aes(y = UK37, x = Year)) +
```

```

geom_point() +
  geom_ribbon(data = newBraya,
    mapping = aes(x = Year, ymax = upper, ymin = lower,
                  fill = Method),
    alpha = 0.3, inherit.aes = FALSE) +
  geom_line(data = newBraya,
    mapping = aes(y = fit, x = Year, colour = Method)) +
  labs(y = braya_ylabel, x = "Year CE") +
  scale_color_manual(values = c("#5e3c99", "#e66101")) +
  scale_fill_manual(values = c("#5e3c99", "#e66101")) +
  theme(legend.position = "right")
braya_fitted

```



133

134 Figure 5 in the manuscript was produced using:

```

plot_grid(small_fitted, braya_fitted, ncol = 1, labels = "auto",
          align = "hv", axis = "lr")

```

135 **3.2.1 Checking if the size of the basis expansion is sufficient**

136 Clearly we are not able to identify both the wiggly trend and the CAR(1) process in a single
 137 model. We might proceed by fitting a simple GAM via `gam()` with REML smoothness selection
 138 and moderate number of basis functions — here we set `k = 15` for illustration.

```

braya_low_k <- gam(UK37 ~ s(Year, k = 15), data = braya, method = "REML")

```

139 Before proceeding, we should perform a check to determine if the number of basis functions re-
 140 quired is sufficient to capture the wiggleness in the data. This test is provided by `gam.check()`

```

gam.check(braya_low_k)

141 #>
142 #> Method: REML   Optimizer: outer newton
143 #> full convergence after 6 iterations.
144 #> Gradient range [-4.468978e-08,8.696688e-10]
145 #> (score -187.2524 & scale 0.000689802).
146 #> Hessian positive definite, eigenvalue range [0.5525272,43.5153].
147 #> Model rank = 15 / 15
148 #>
149 #> Basis dimension (k) checking results. Low p-value (k-index<1) may
150 #> indicate that k is too low, especially if edf is close to k'.
151 #>
152 #>          k'    edf k-index p-value
153 #> s(Year) 14.00  2.62    0.56  <2e-16 ***
154 #> ---
155 #> Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

156 In the output, the column labelled k' is the size of the basis used (the number of basis functions)
157 — note that it is 14 and not the requested 15 functions as the constant basis function has been
158 removed as it is confounded with the model constant term, the intercept. The column labelled
159 edf indicates the effective degrees of freedom for the estimated smooth. The k-index column
160 contains the value of a test statistic for a test of sufficient number of basis functions — ideally,
161 we'd like the value of k-index to be close to or greater than 1. The quoted p value measures the
162 support for the null hypothesis that enough basis functions were available. In practical terms,
163 the low k-index and very low p value suggest that the requested number of basis functions
164 used to fit the model was too low.

165 The solution is to increase k, say by doubling the original value, and refit the model and per-
166 form the basis dimension check using gam.check(), repeating this process until k-index is
167 close to or greater than 1 and the p value is larger than a desired threshold. In my experience,
168 it may not be possible with some data sets to increase k large enough to result in a k-index > 1
169 and a p value > 0.05, given the available. It is important to remember that the test provided by
170 gam.check() is only an heuristic one and shouldn't be seen as infallible. If you find yourself
171 increasing k to large values relative to the number of samples in your data set without achiev-
172 ing a k-index > 1 or p > 0.05, then note the value in the edf column as you increase k. If the
173 edf of the model changes little as you continue to increase k, this may be evidence that the
174 test is failing in this particular instance and that the estimated smooth is not dependent on the
175 value of k chosen.

```

176 3.3 Accounting for heteroscedasticity due to time averaging

177 To proceed with the Braya-Sø example, we need to increase the basis dimension ($k = 45$),
178 fit using `method = "REML"`. As discussed in the manuscript, we should also account for the
179 non-constant variance, or *heteroscedasticity*, that may arise due to variation in the amount of

180 time (or “lake years”) that each sample represents. All else equal, we expect that samples
 181 that average more time have smaller variance than those that average a smaller amount of
 182 time. To include information about the expected heteroscedasticity in the GAM we can use
 183 observational weights². Here I use the `sampleInterval` variable that I created earlier as the
 184 measure of lake years per sample, and to avoid changing the model likelihood, the weights
 185 use are the values of `sampleInterval` divided by the mean of `sampleInterval`:

```
braya_reml <- gam(UK37 ~ s(Year, k = 45), data = braya,
  method = "REML",
  weights = sampleInterval / mean(sampleInterval))
summary(braya_reml)
```

```
186 #>
187 #> Family: gaussian
188 #> Link function: identity
189 #>
190 #> Formula:
191 #> UK37 ~ s(Year, k = 45)
192 #>
193 #> Parametric coefficients:
194 #>             Estimate Std. Error t value Pr(>|t|)
195 #> (Intercept) -0.633734  0.001901 -333.3   <2e-16 ***
196 #> ---
197 #> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
198 #>
199 #> Approximate significance of smooth terms:
200 #>             edf Ref.df    F p-value
201 #> s(Year) 28.49 33.99 7.755 <2e-16 ***
202 #> ---
203 #> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
204 #>
205 #> R-sq.(adj) =  0.751  Deviance explained = 83.2%
206 #> -REML = -175.68  Scale est. = 0.00024036 n = 89
207 If we now check if 44 (k = 45) basis functions is sufficient
```

```
gam.check(braya_reml)
```

```
208 #>
209 #> Method: REML Optimizer: outer newton
210 #> full convergence after 7 iterations.
211 #> Gradient range [-2.210576e-06,-5.181457e-08]
212 #> (score -175.684 & scale 0.0002403604).
213 #> Hessian positive definite, eigenvalue range [2.524305,48.11146].
```

²This simple solution is for a Gaussian GAM. For models with other distributions, the exact interpretation of observational weights will vary. For such data, where the conditional mean and variance of the observations are linked, the use of location-scale, or distributional, models may be more appropriate and informative.

```

214 #> Model rank = 45 / 45
215 #>
216 #> Basis dimension (k) checking results. Low p-value (k-index<1) may
217 #> indicate that k is too low, especially if edf is close to k'.
218 #>
219 #>          k'   edf k-index p-value
220 #> s(Year) 44.0 28.5    1.27    0.99
221 we note that the value of k-index is now greater than 1 and the p value suggests there is very
222 little evidence against the null hypothesis. Also note that the edf of the estimated smooth is
223 considerably below the maximum (k'). This is the effect of the smoothness penalty removing
224 the excessive wigginess possible with 44 basis functions. If the edf were close to k', you
225 might consider increasing k by a modest amount (in this instance by 5 or 10 additional basis
226 functions) to assure yourself that you have sufficient basis functions. This has the additional
227 advantage of allowing a richer set of smooths of a given complexity (edf) to be represented
228 using the basis functions, which may improve the model fit to the data without changing the
229 edf too much.

```

230 4 Posterior simulation

231 Samples from the posterior distribution of a GAM can be drawn using the `simulate()` meth-
232 ods from the *schoenberg* package.

```


set.seed(1) # set the random seed to make this reproducible
nsim <- 20 # how many simulations to draw

## do the simulations
sims <- simulate(mod, nsim = nsim, newdata = newYear, unconditional = TRUE)

## rearrange the output into a long/tidy format
colnames(sims) <- paste0("sim", seq_len(nsim))
sims <- setNames(stack(as.data.frame(sims)), c("simulated", "run"))
sims <- transform(sims, Year = rep(newYear$Year, nsim),
                  simulated = simulated)

## Plot simulated trends
smallSim.plt <- ggplot(newYear, aes(x = Year, y = fit)) +
  geom_line(data = sims,
             mapping = aes(y = simulated, x = Year, group = run),
             colour = "grey80") +
  geom_line(lwd = 1) +
  labs(y = d15n_label, x = "Year CE")
smallSim.plt

```



233

234 We repeat the same simulation for Braya-Sø

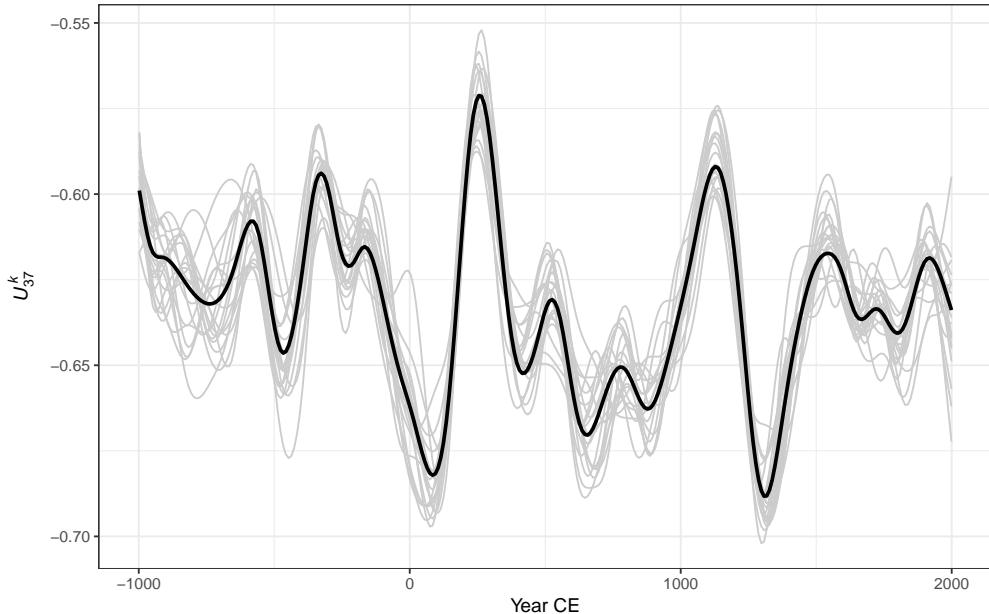
```

## posterior simulation
## need to reset-up newBraya
newBraya <- with(braya,
                 data.frame(Year = seq(min(Year), max(Year),
                                         length.out = N)))
braya_pred <- cbind(newBraya,
                      data.frame(predict(braya_reml, newBraya,
                                         se.fit = TRUE)))

## simulate
set.seed(1)
sims2 <- simulate(braya_reml, nsim = nsim, newdata = newBraya,
                    unconditional = TRUE)
colnames(sims2) <- paste0("sim", seq_len(nsim))
sims2 <- setNames(stack(as.data.frame(sims2)),
                  c("simulated", "run"))
sims2 <- transform(sims2, Year = rep(newBraya$Year, nsim),
                    simulated = simulated)

brayaSim.plt <- ggplot(braya_pred, aes(x = Year, y = fit)) +
  geom_line(data = sims2,
            mapping = aes(y = simulated, x = Year, group = run),
            colour = "grey80") +
  geom_line(lwd = 1) +
  labs(y = braya_ylabel, x = "Year CE")
brayaSim.plt

```



235

236 Figure 7 in the manuscript was prepared using

```
plot_grid(smallSim.plt, brayaSim.plt, ncol = 1, labels = "auto",
          align = "hv", axis = "lr")
```

237 5 Confidence and simultaneous intervals

238 Across-the-function and simultaneous confidence intervals are computed using the
 239 `confint()` method. The type of interval required is given via the `type` argument with
 240 options "confidence" and "simultaneous". For example, for Small Water we would use

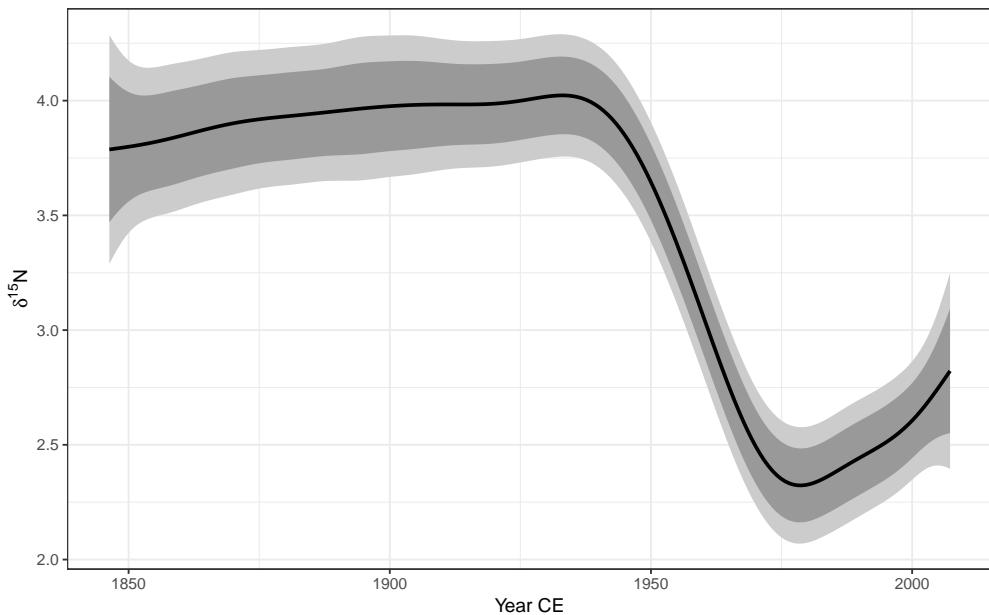
```
sw.cint <- confint(mod, parm = "Year", newdata = newYear,
                     type = "confidence")
sw.sint <- confint(mod, parm = "Year", newdata = newYear,
                     type = "simultaneous")
head(sw.sint)
```

```
241 #>   smooth     Year      est       se      crit    lower     upper
242 #> 1 s(Year) 1846.319 3.787215 0.1616144 3.087688 3.288200 4.286230
243 #> 2 s(Year) 1847.128 3.789747 0.1509100 3.087688 3.323784 4.255710
244 #> 3 s(Year) 1847.937 3.792304 0.14111439 3.087688 3.356496 4.228112
245 #> 4 s(Year) 1848.745 3.794914 0.1324961 3.087688 3.385807 4.204020
246 #> 5 s(Year) 1849.554 3.797602 0.1251070 3.087688 3.411310 4.183893
247 #> 6 s(Year) 1850.363 3.800394 0.1190552 3.087688 3.432788 4.167999
```

248 The `confint()` methods return data frames suitable for subsequent plotting with **ggplot**. The
 249 columns labelled `est` and `se` are the estimate values of the smooth and its standard error,

250 respectively. The variables lower and upper contain the values of the lower and upper bounds
251 on the requested interval. The intervals can be plotted as follows

```
smallInt.plt <- ggplot(sw.cint, aes(x = Year, y = est)) +  
  geom_ribbon(data = sw.sint,  
              mapping = aes(ymin = lower, ymax = upper, x = Year),  
              fill = "grey80", inherit.aes = FALSE) +  
  geom_ribbon(mapping = aes(ymin = lower, ymax = upper, x = Year),  
              fill = "grey60", inherit.aes = FALSE) +  
  geom_line(lwd = 1) +  
  labs(y = d15n_label, x = "Year CE")  
smallInt.plt
```



252

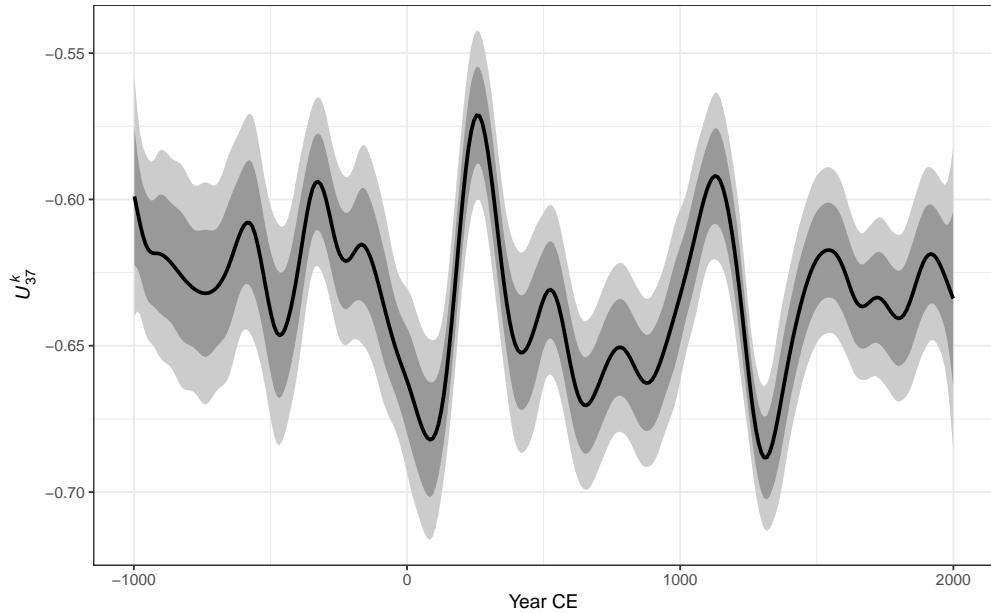
253 The intervals for Braya-Sø are created

```
bs.cint <- confint(braya_reml, parm = "Year", newdata = newBraya,  
                     type = "confidence")  
bs.sint <- confint(braya_reml, parm = "Year", newdata = newBraya,  
                     type = "simultaneous")
```

254 and plotted

```
brayaInt.plt <- ggplot(bs.cint, aes(x = Year, y = est)) +  
  geom_ribbon(data = bs.sint,  
              mapping = aes(ymin = lower, ymax = upper, x = Year),  
              fill = "grey80", inherit.aes = FALSE) +  
  geom_ribbon(mapping = aes(ymin = lower, ymax = upper, x = Year),  
              fill = "grey60", inherit.aes = FALSE) +  
  geom_line(lwd = 1) +  
  labs(y = braya_ylabel, x = "Year CE")
```

```
brayaInt.plt
```



255

256 in the same way

257 Figure 8 in the manuscript was prepared using

```
plot_grid(smallInt.plt, brayaInt.plt, ncol = 1, labels = "auto",
          align = "hv", axis = "lr")
```

258 6 Derivatives of the estimated trend

259 The first derivative of the estimated trend is calculated using finite differences using the
260 `fderiv()` function. There is also a `confint()` method for objects produced by `fderiv()`. The
261 first derivatives and a 95% simultaneous confidence interval for the Small Water trend we
262 computed and plotted using

```
small.d <- fderiv(mod, newdata = newYear, n = N)
small.sint <- with(newYear,
                    cbind(confint(small.d, nsim = nsim,
                                  type = "simultaneous"),
                          Year = Year))

small_deriv_plt <- ggplot(small.sint, aes(x = Year, y = est)) +
  geom_ribbon(aes(ymin = lower, ymax = upper), alpha = 0.2,
              fill = "black") +
  geom_line() +
  labs(x = "Year CE", y = "First derivative")
```

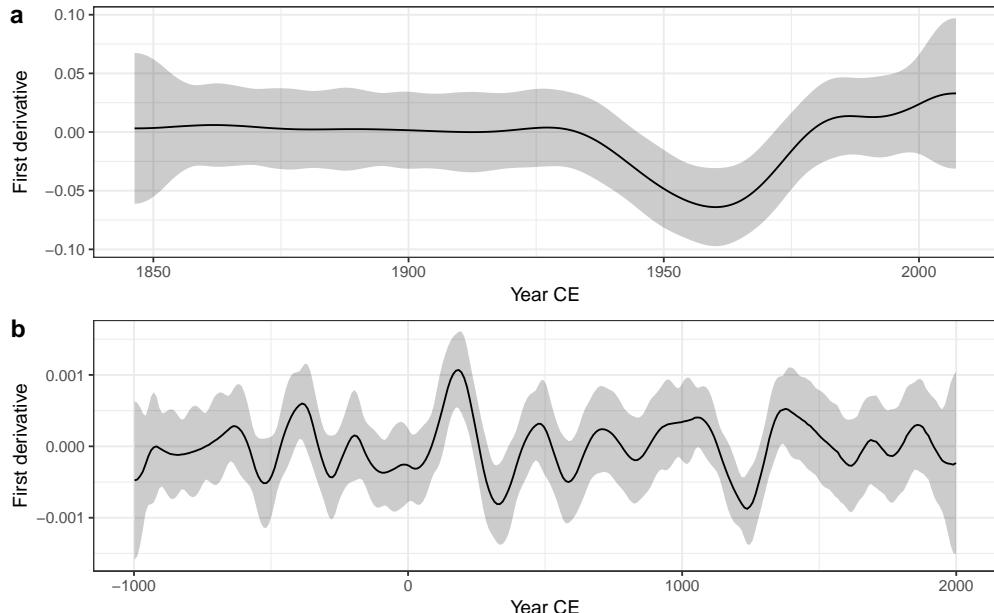
263 whilst for Braya-Sø, the following was used

```
braya.d <- fderiv(braya_reml, newdata = newBraya, n = N)
braya.sint <- with(newBraya,
  cbind(confint(braya.d, nsim = nsim,
    type = "simultaneous"),
  Year = Year))

braya_deriv_plt <- ggplot(braya.sint, aes(x = Year, y = est)) +
  geom_ribbon(aes(ymin = lower, ymax = upper),
    alpha = 0.2, fill = "black") +
  geom_line() +
  labs(x = "Year CE", y = "First derivative")
```

264 Figure 9 in the manuscript was prepared using

```
plot_grid(small_deriv_plt, braya_deriv_plt, ncol = 1, labels = "auto",
  align = "hv", axis = "lr")
```



265

266 7 Gaussian process smooths

267 For the Gaussian process smooth to fit within the GAM framework described in the
268 manuscript, we need to supply the value of ϕ for the effective range of the correlation function.
269 To estimate ϕ , we need to repeatedly fit the required GAM using a range of plausible values
270 for ϕ , which we do using a loop. In the chunk below I fit 200 models with ϕ in the range
271 15–500. For value of ϕ I fit the required GAM and extract the REML score (from component
272 `gcv.ubre`) and store it in the numeric vectors `Mat` or `SEx`, for Matérn and Squared Exponential

273 correlation functions, respectively. The final line of the chunk prepares the REML scores for
274 the two correlation functions in long or tidy format suitable for plotting with *ggplot2*.

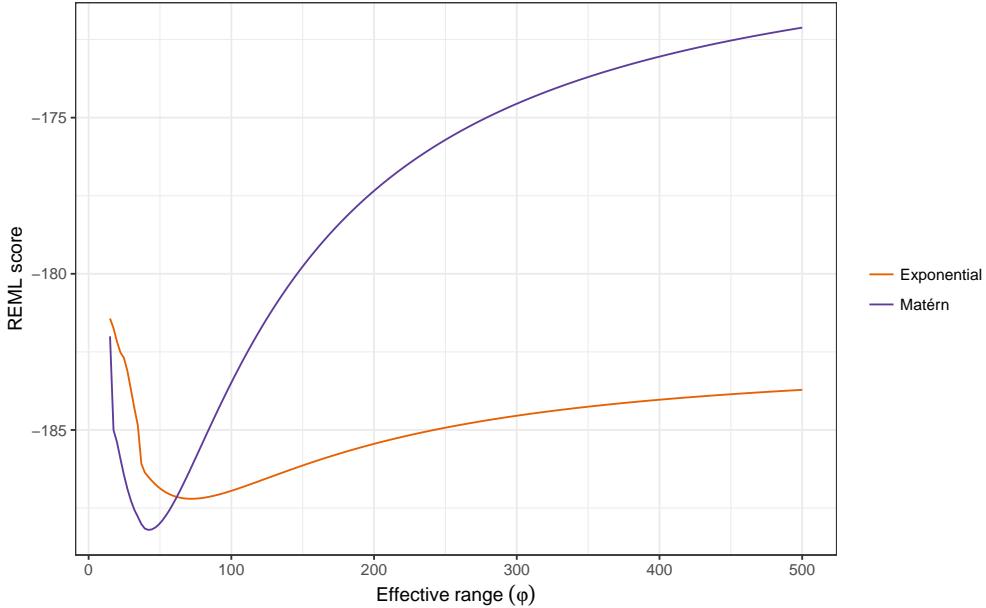
```
nn <- 200      # number of points at which to evaluate profile likelihood
dseq <- seq(15, 500, length.out = nn)  # effective ranges to fit at
Mat <- SEx <- numeric(length = nn)      # object to hold model fits

## loop over the sequence of separation distances and fit the models
for (i in seq_along(dseq)) {
    ## iterate over dseq, fit GP GAM w Matérn covariance
    Mat[i] <- gam(UK37 ~ s(Year, k = 45, bs = "gp", m = c(3, dseq[i])),
                   weights = sampleInterval / mean(sampleInterval),
                   data = braya, method = "REML",
                   family = gaussian())[["gcv.ubre"]]
    ## fit squared exponential
    SEx[i] <- gam(UK37 ~ s(Year, k = 45, bs = "gp", m = c(2, dseq[i], 1)),
                   weights = sampleInterval / mean(sampleInterval),
                   data = braya, method = "REML",
                   family = gaussian())[["gcv.ubre"]]
}

## extract the REML score into ggplot-friendly object
reml.scr <- data.frame(cor = rep(c("Matérn", "Exponential"), each = nn),
                        effrange = rep(dseq, 2),
                        reml = c(Mat, SEx))
```

275 The REML scores for the models are plotted with *ggplot2* using

```
## profile-likelihood plot
proflik.plt <- ggplot(reml.scr, aes(x = effrange, y = reml, colour = cor)) +
    geom_line() +
    scale_colour_manual(name = "", values = c("#e66101", "#5e3c99")) +
    labs(y = "REML score", x = expression(Effective ~ range ~ (varphi)))
proflik.plt
```



276

277 Next we extract the minimum of the REML scores for the two correlation functions and refit
 278 those models (we threw away all the models in the `for ()` loop earlier to avoid storing lots
 279 of model objects). Then we fit GAMs with Gaussian process smooths using the values of ϕ
 280 that produced the minimum REML scores, and predict using the fitted models to visualize
 281 the trends.

```
## effective range minima from profile likelihood
effRange2 <- with(subset(reml.scr, cor == "Matérn"), dseq[which.min(reml)])
effRange3 <- with(subset(reml.scr, cor == "Exponential"), dseq[which.min(reml)])

## Refit these models: Matern
gp2 <- gam(UK37 ~ s(Year, k = 45, bs = "gp", m = c(3, effRange2)),
            data = braya,
            method = "REML", weights = sampleInterval / mean(sampleInterval))
## Refit these models: Power exponential
gp3 <- gam(UK37 ~ s(Year, k = 45, bs = "gp", m = c(2, effRange3, 1)),
            data = braya,
            method = "REML", weights = sampleInterval / mean(sampleInterval))

## data to predict at
newd <- with(braya, data.frame(Year = seq(min(Year), max(Year),
                                         length.out = 1000)))
## create predictions on response scale for both covariance functions
p.gp2 <- transform(newd,
                     fitted = predict(gp2, newdata = newd, type = "response"),
                     effRange = round(effRange2))
p.gp3 <- transform(newd,
                     fitted = predict(gp3, newdata = newd, type = "response"),
```

```

    effRange = round(effRange3))
## joint the two sets of predictions together
pred <- rbind(p.gp2, p.gp3)
## add some categorical variables for plotting
pred <- transform(pred,
                   effRange = factor(effRange),
                   cor = rep(c("Matérn", "Exponential"), each = nrow(newd)))

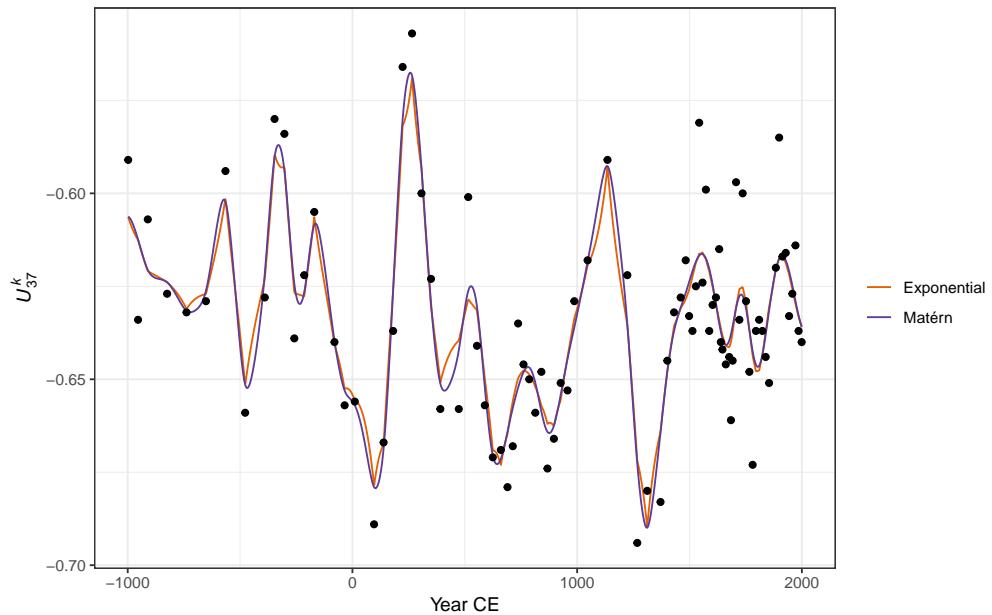
```

- 282 The estimated Gaussian process trends are plotted using

```

gp.plt2 <- ggplot(pred, aes(x = Year, y = fitted, colour = cor)) +
  geom_line() + theme(legend.position = "right") +
  geom_point(aes(x = Year, y = UK37), data = braya, inherit.aes = FALSE) +
  scale_colour_manual(name = "", values = c("#e66101", "#5e3c99")) +
  labs(y = braya_ylabel, x = "Year CE")
gp.plt2

```



283

- 284 whilst Figure 12 in the manuscript was prepared using

```

plot_grid(proflik.plt, gp.plt2, ncol = 1, labels = c("a", "b"),
          align = "hv", axis = "lr")

```

285 8 Adaptive smooths

- 286 The adaptive smooth was fitted to the Braya-Sø data by adding `bs = "ad"` to the `s()` term in
 287 the model formula. The other aspects of the fit are as previously used for the other models,
 288 REML smoothness selection and observational weights:

```

## Adaptive spline, weights as sampleInterval
mod_ad <- gam(UK37 ~ s(Year, k = 45, bs = "ad"), data = braya,
               method = "REML",
               weights = sampleInterval / mean(sampleInterval))

```

289 9 Comparing trends

- 290 For the model comparison, I refitted all the models for consistency; the code to fit each of the
 291 1. Thin plate regression spline (TPRS),
 292 2. Gaussian process spline (Matérn correlation functions), and
 293 3. Adaptive smoother

294 is shown below.

```

## effective range parameter from profile-likelihood of Matern model
effRange <- effRange2

## TPRS, weights as sampleInterval
mod_tprs <- gam(UK37 ~ s(Year, k = 45, bs = "tp"), data = braya,
                  method = "REML",
                  weights = sampleInterval / mean(sampleInterval))

## Gaussian process, Matern, kappa = 1.5, weights as sampleInterval
mod_gp <- gam(UK37 ~ s(Year, k = 45, bs = "gp", m = c(3, effRange)),
                data = braya,
                method = "REML",
                weights = sampleInterval / mean(sampleInterval))

## Adaptive spline, weights as sampleInterval
mod_ad <- gam(UK37 ~ s(Year, k = 45, bs = "ad"), data = braya,
                 method = "REML",
                 weights = sampleInterval / mean(sampleInterval))

```

- 295 We write a small function to predict from each model over the range of Year and return the
 296 data in tidy format for plotting. The plot produce reproduces figure 13 in the manuscript.

```

## wrap this in a function that will return all the plots & derived objects
processGAM <- function(mod) {
  ## Predict from model
  N <- 500
  newYear <- with(braya,
                  data.frame(Year = seq(min(Year), max(Year),
                                         length.out = N)))
  newYear <- cbind(newYear,

```

```

        data.frame(predict(mod, newYear, se.fit = TRUE)))


out <- list(objects = newYear)
out
}

plts_gp   <- processGAM(mod = mod_gp) # Gaussian process smooth with weights
plts_ad   <- processGAM(mod = mod_ad) # Adaptive smooth with weights
plts_tprs <- processGAM(mod = mod_tprs) # TPRS with weights

pltData <- do.call("rbind", lapply(list(plts_gp, plts_ad, plts_tprs),
                                     `[[`, "objects")))
pltData <- transform=pltData, Model = rep(c("GP", "Adaptive", "TPRS"),
                                         each = nrow(plts_gp$objects)))

allFits <- ggplot=pltData, aes(x = Year, y = fit)) +
  geom_point(aes(x = Year, y = UK37), data = braya) +
  geom_line(aes(colour = Model)) + labs(y = braya_ylabel, x = "Year") +
  theme(legend.position = "right") +
  scale_colour_manual(name = "",
                      values = c("#e66101", "#fdb863", "#5e3c99"))
allFits

```



298 10 Accounting for age-model uncertainty

299 The manuscript proposed to simulate from the posterior distribution of the fitted age model
300 as a way to account for age-model uncertainty. The first step in the process is to fit the age
301 model from which to simulate new age models. This was done using the *scam* package for a
302 *shape-constrained GAM*, with the age-model spline constrained to be monotonic decreasing (*bs*
303 = "mpd").

304 To make this section self-contained, I refitted the Small Water GAM plus CAR(1) model

```
knots <- with(small, list(Year = seq(min(Year), max(Year), length = 14)))  
mod <- gamm(d15N ~ s(Year, k = 15), data = small, method = "REML",  
            correlation = corCAR1(form = ~ Year),  
            knots = knots)
```

305 and then we load the ^{210}Pb dating results for the dated core sections.

```
swAge <- read.csv("./data/small-water/small11-dating.csv")
```

306 before fitting the shape-constrained GAM. Currently, *scam* can only fit models using GCV
307 smoothness selection. I used the *gamma* argument here to add a larger penalty for more-
308 complex models. Each effective degree of freedom used by the spline is counted as 1.4
309 degrees of freedom in the GCV score.

```
## monotonic spline age-depth model  
swAge$Error[1] <- 1.1  
swAgeMod <- scam(Date ~ s(Depth, k = 5, bs = "mpd"), data = swAge,  
                  weights = 1 / swAge$Error, gamma = 1.4)
```

310 Note that I added a small amount of error to the surface sample age as the model cannot be
311 fitted if an observation has 0 weight.

312 Next, predict from the estimated age model, and draw 25 samples from the posterior distribution
313 using *simulate()*. The results are tidied into a format suitable for further processing and
314 plotting. Note that the posterior samples here are only used for plotting.

```
## predict from the age model for a smooth set of points in `Depth`  
newAge <- with(swAge, data.frame(Depth = seq(min(Depth), max(Depth),  
                                    length.out = 200)))  
newAge <- transform(newAge,  
                     fitted = predict(swAgeMod, newdata = newAge,  
                                       type = "response"))  
newSims <- as.data.frame(simulate(swAgeMod, nsim = 25, newdata = newAge))  
newSims <- cbind(Depth = newAge$Depth, newSims)  
newSims <- gather(newSims, Simulation, Age, -Depth)
```

315 In the next code chunk, I draw 100 samples from the posterior distribution of the age model,
316 but notice that I pass in the *small* data to *newdata* in the call to *simulate()* as the locations
317 I want new age estimates for are the depths for which we have $\delta^{15}\text{N}$ values. A small function

318 (fitSWModels) is written to prepare each simulation for fitting and then actually fit the GAM
319 plus CAR(1) model using the updated age information.

```
## simulate from age model; each column is a simulation
ageSims <- simulate(swAgeMod, nsim = 100, newdata = small, seed = 42)
ageSims <- as.data.frame(ageSims)

fitSWModels <- function(x, y, knots) {
  dat <- data.frame(d15N = y, Year = x)
  m <- gamm(d15N ~ s(Year, k = 15), data = dat, method = "REML",
             correlation = corCAR1(form = ~ Year), knots = knots)
}

## generate new trends using draws from age-model posterior
simTrendMods <- lapply(ageSims, fitSWModels, y = small$d15N, knots = knots)

## function wrapper to predict new trends at locations over the
## range of `Year`
predSWModels <- function(mod, newdata) {
  predict(mod$gam, newdata = newdata, type = "response")
}

## predict from fitted model to produce a smooth trend for each posterior
## sample
simTrends <- lapply(simTrendMods, predSWModels, newdata = newYear)

## arrange in a tidy format for plottings
simTrends <- data.frame(Year = with(newYear, rep(Year, length(simTrends))),
                        Trend = unlist(simTrends),
                        Group = rep(seq_along(simTrends),
                                   times = lengths(simTrends)))
```

320 The next chunk does the final step in the process. For each of the models we just fitted to
321 include age model uncertainty, we simulate 50 draws from the model posterior distribution.
322 We start with a wrapper function around the simulate() code we want to run on each model,
323 then do the actual posterior draws for each model using lapply(). The final step just arranges
324 data for plotting.

```
## wrapper to simulate from a fitted GAM with the required arguments
simulateSWModels <- function(mod, newdata, nsim, seed = 42) {
  sims <- simulate(mod, nsim = nsim, newdata = newdata, seed = seed)
  as.vector(sims)
}

## now do the posterior simulation
```

```

NSIM <- 50      # number of posterior samples *per* model
simSimulate <- lapply(simTrendMods, simulateSWModels, newdata = newYear,
                      nsim = NSIM, seed = 42)

## arrange in a tidy format
simSimulate <-
  data.frame(Year = with(newYear,
                         rep(Year, times = NSIM * length(simSimulate))),
             Trend = unlist(simSimulate),
             Group = rep(seq_len(NSIM * length(simSimulate)),
                         each = nrow(newYear)))

```

- 325 Each of the steps is visualized using the plot code shown below.

```

plt1 <- ggplot(swAge, aes(y = Date, x = Depth)) +
  geom_line(data = newSims,
            mapping = aes(y = Age, x = Depth, group = Simulation),
            alpha = 1, colour = "grey80") +
  geom_line(data = newAge, mapping = aes(y = fitted, x = Depth)) +
  geom_point(size = 1.5, colour = "red") +
  geom_errorbar(aes(ymin = Date - Error, ymax = Date + Error, width = 0),
                colour = "red") +
  labs(y = "Year CE", x = "Depth (cm)")

plt2 <- ggplot(simTrends, aes(x = Year, y = Trend, group = Group)) +
  geom_line(alpha = 0.1, colour = "grey80") +
  geom_line(data = newYear,
            mapping = aes(x = Year, y = fit), inherit.aes = FALSE) +
  geom_point(data = small,
             mapping = aes(x = Year, y = d15N),
             inherit.aes = FALSE, size = 0.7) +
  labs(x = "Year", y = d15n_label)

plt3 <- ggplot(simSimulate, aes(x = Year, y = Trend, group = Group)) +
  geom_line(alpha = 0.2, colour = "grey80") +
  geom_point(data = small,
             mapping = aes(x = Year, y = d15N),
             inherit.aes = FALSE,
             size = 0.7) +
  geom_line(data = newYear,
            mapping = aes(x = Year, y = fit),
            inherit.aes = FALSE) +
  labs(x = "Year", y = d15n_label)

## align all plots vertically

```

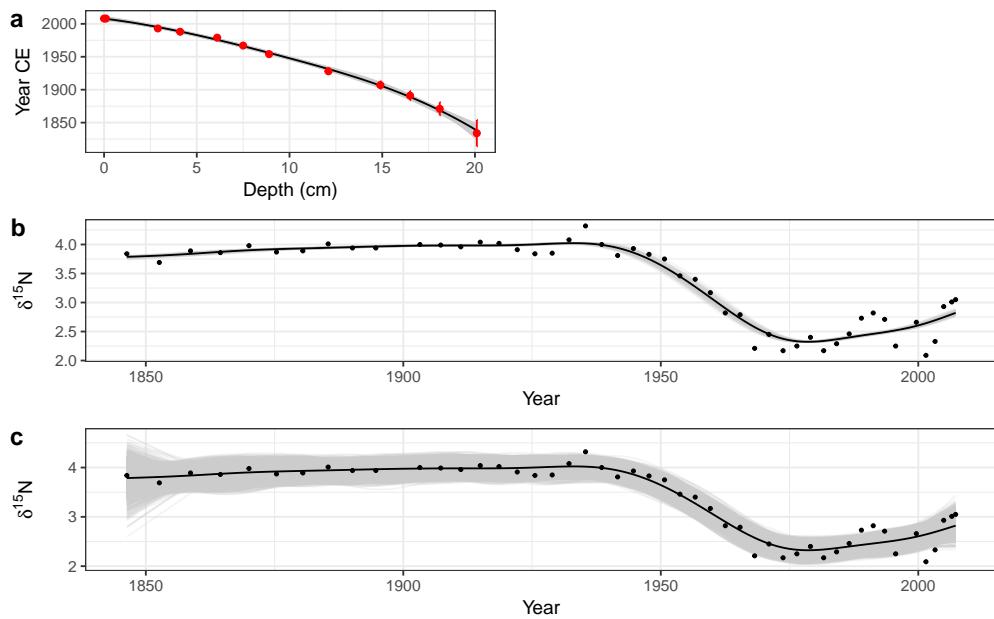
```

plots <- align_plots(plt1, plt2, plt3, align = 'v', axis = 'l')

## create the two rows of figures, from `plots`
top_row <- plot_grid(plots[[1]], NULL, ncol = 2, labels = "a")
bot_row <- plot_grid(plots[[2]], plots[[3]], ncol = 1, labels = c("b", "c"))

## combine the two rows, top row has 1 plot row, bottom row has 2, hence
## the rel_heights to even this out
plot_grid(top_row, bot_row, ncol = 1, rel_heights = c(0.5, 1))

```



326

327 This reproduces figure 14 from the manuscript.

328 11 Session information

```
devtools::session_info()
```

```

329 #> Session info -----
330 #>   setting  value
331 #>   version R version 3.4.4 Patched (2018-03-17 r74422)
332 #>   system  x86_64, linux-gnu
333 #>   ui      X11
334 #>   language (EN)
335 #>   collate en_CA.UTF-8
336 #>   tz      America/Regina

```

```

337 #> date      2018-05-14
338 #> Packages -----
339 #> package    * version date      source
340 #> backports   1.1.2   2017-12-13 CRAN (R 3.4.3)
341 #> base        * 3.4.4   2018-03-18 local
342 #> codetools    0.2-15  2016-10-05 CRAN (R 3.4.4)
343 #> colorspace   1.3-2   2016-12-14 CRAN (R 3.4.1)
344 #> compiler     3.4.4   2018-03-18 local
345 #> cowplot      * 0.9.2  2017-12-17 CRAN (R 3.4.3)
346 #> datasets     * 3.4.4   2018-03-18 local
347 #> devtools     1.13.5  2018-02-18 CRAN (R 3.4.4)
348 #> digest       0.6.15  2018-01-28 CRAN (R 3.4.4)
349 #> evaluate     0.10.1  2017-06-24 CRAN (R 3.4.1)
350 #> ggplot2      * 2.2.1  2016-12-30 CRAN (R 3.4.1)
351 #> glue         1.2.0   2017-10-29 CRAN (R 3.4.2)
352 #> graphics     * 3.4.4   2018-03-18 local
353 #> grDevices    * 3.4.4   2018-03-18 local
354 #> grid          3.4.4   2018-03-18 local
355 #> gtable        0.2.0   2016-02-26 CRAN (R 3.4.1)
356 #> htmltools     0.3.6   2017-04-28 CRAN (R 3.4.1)
357 #> knitr         1.20    2018-02-20 CRAN (R 3.4.4)
358 #> labeling      0.3     2014-08-23 CRAN (R 3.4.1)
359 #> lattice        0.20-35 2017-03-25 CRAN (R 3.4.4)
360 #> lazyeval      0.2.1   2017-10-29 CRAN (R 3.4.2)
361 #> magrittr      1.5     2014-11-22 CRAN (R 3.4.1)
362 #> MASS          7.3-49  2018-02-23 CRAN (R 3.4.4)
363 #> Matrix        1.2-14  2018-04-09 CRAN (R 3.4.4)
364 #> memoise       1.1.0   2017-04-21 CRAN (R 3.4.1)
365 #> methods       3.4.4   2018-03-18 local
366 #> mgcv          * 1.8-23 2018-01-21 CRAN (R 3.4.4)
367 #> munsell       0.4.3   2016-02-13 CRAN (R 3.4.1)
368 #> nlme          * 3.1-137 2018-04-07 CRAN (R 3.4.4)
369 #> pillar         1.2.1   2018-02-27 CRAN (R 3.4.4)
370 #> plyr          1.8.4   2016-06-08 CRAN (R 3.4.1)
371 #> purrr         0.2.4   2017-10-18 CRAN (R 3.4.2)
372 #> Rcpp          0.12.16 2018-03-13 CRAN (R 3.4.4)
373 #> rlang          0.2.0   2018-02-20 CRAN (R 3.4.4)
374 #> rmarkdown     * 1.9     2018-03-01 CRAN (R 3.4.4)
375 #> rprojroot     1.3-2   2018-01-03 CRAN (R 3.4.3)
376 #> scales         0.5.0   2017-08-24 CRAN (R 3.4.1)
377 #> scam           * 1.2-2   2017-09-24 CRAN (R 3.4.3)
378 #> schoenberg   * 0.0-6   2018-04-12 Github (gavinsimpson/schoenberg@f9b8a84)
379 #> splines        3.4.4   2018-03-18 local
380 #> stats          * 3.4.4   2018-03-18 local

```

```
381 #> stringi      1.1.7  2018-03-12 CRAN (R 3.4.4)
382 #> stringr      1.3.0  2018-02-19 CRAN (R 3.4.4)
383 #> tibble        1.4.2  2018-01-22 CRAN (R 3.4.3)
384 #> tidyverse     * 0.8.0  2018-01-29 CRAN (R 3.4.4)
385 #> tools         3.4.4  2018-03-18 local
386 #> utils         * 3.4.4  2018-03-18 local
387 #> withr        2.1.2  2018-03-15 CRAN (R 3.4.4)
388 #> yaml          2.1.18 2018-03-08 CRAN (R 3.4.4)
```